

## 1. Simple test for each function

Test multiply\_mv\_row\_major (should be [6, 15]):

Result of multiply\_mv\_row\_major:

6 15

Test multiply\_mv\_row\_major (should be [6, 15]):

Result of multiply\_mv\_col\_major:

6 15

Test multiply\_mm\_naive (should be [58 64; 139 154]):

Result of multiply\_mm\_naive:

58 64

139 154

Test multiply\_mm\_naive (should be [58 64; 139 154]):

Result of multiply\_mm\_transposed\_b:

58 64

139 154

## 2. Benchmarking

Size: 3

Matrix-Vector-Row | Avg Time: 8.01e-05 ms | Std Dev: 2.48397e-05 ms

Matrix-Vector-Column | Avg Time: 9.25e-05 ms | Std Dev: 2.39502e-05 ms

Size: 3x3

Matrix-Matrix-Naive | Avg Time: 0.00018424 ms | Std Dev: 2.76879e-05 ms

Matrix-Matrix-Transposed | Avg Time: 0.00017916 ms | Std Dev: 2.23842e-05 ms

Size: 20

Matrix-Vector-Row | Avg Time: 0.00221586 ms | Std Dev: 0.000907416 ms

Matrix-Vector-Column | Avg Time: 0.00220914 ms | Std Dev: 0.000973124 ms

Size: 20x20

Matrix-Matrix-Naive | Avg Time: 0.0390959 ms | Std Dev: 0.00811997 ms

Matrix-Matrix-Transposed | Avg Time: 0.0382925 ms | Std Dev: 0.00705416 ms

Size: 100

Matrix-Vector-Row | Avg Time: 0.0491375 ms | Std Dev: 0.00783763 ms

Matrix-Vector-Column | Avg Time: 0.0486659 ms | Std Dev: 0.00699852 ms

Size: 100x100

Matrix-Matrix-Naive | Avg Time: 3.4227 ms | Std Dev: 0.15749 ms

Matrix-Matrix-Transposed | Avg Time: 3.41261 ms | Std Dev: 0.124625 ms

### 3. Cache Locality Analysis

Matrix-vector multiplication (row-major vs. column-major):

In the row-major implementation, memory is accessed sequentially along rows, which aligns with CPU cache line behavior and takes advantage of spatial locality. In contrast, the column-major implementation accesses memory non-sequentially along columns, leading to more cache misses due to poor spatial locality. Therefore, the row-major version performs better in terms of cache efficiency.

Matrix-matrix multiplication (naive vs. transposed):

The naive implementation accesses matrix B in a column-wise fashion, which results in non-contiguous memory access when B is stored in row-major order. This leads to frequent cache misses. In the transposed B implementation, matrix B is first transposed, enabling row-wise access and therefore sequential memory access. This improves spatial locality and reduces cache misses. As a result, the transposed B version is expected to perform better, especially for larger matrices.

### 4. Memory Alignment

```
==== Unaligned Benchmark ====
Size: 3x3
Matrix-Matrix-Naive | Avg Time: 8.912e-05 ms | Std Dev: 1.67423e-05 ms
Matrix-Matrix-Transposed | Avg Time: 9.336e-05 ms | Std Dev: 1.77851e-05 ms

==== Aligned ====
Size: 3x3
Matrix-Matrix-Naive | Avg Time: 0.0001059 ms | Std Dev: 2.06971e-05 ms
Matrix-Matrix-Transposed | Avg Time: 0.00011252 ms | Std Dev: 2.0754e-05 ms

==== Unaligned Benchmark ====
Size: 20x20
Matrix-Matrix-Naive | Avg Time: 0.0220833 ms | Std Dev: 0.000776644 ms
Matrix-Matrix-Transposed | Avg Time: 0.0219934 ms | Std Dev: 0.0006143 ms

==== Aligned ====
Size: 20x20
Matrix-Matrix-Naive | Avg Time: 0.0219709 ms | Std Dev: 0.000577951 ms
Matrix-Matrix-Transposed | Avg Time: 0.0219692 ms | Std Dev: 0.00044212 ms

==== Unaligned Benchmark ====
Size: 100x100
Matrix-Matrix-Naive | Avg Time: 3.39394 ms | Std Dev: 0.0344013 ms
Matrix-Matrix-Transposed | Avg Time: 3.38375 ms | Std Dev: 0.0423679 ms

==== Aligned ====
Size: 100x100
Matrix-Matrix-Naive | Avg Time: 3.41867 ms | Std Dev: 0.0672101 ms
Matrix-Matrix-Transposed | Avg Time: 3.41993 ms | Std Dev: 0.0637012 ms
```

Memory alignment had little to no noticeable impact on the performance of the matrix operations across all tested sizes. For both the naive and transposed implementations, the average execution times and standard deviations of aligned and unaligned versions were nearly identical.

For small matrices (e.g., 3x3), alignment showed no consistent benefit. For medium (20x20) and large matrices (100x100), the aligned versions did not outperform unaligned ones in any significant way. In some cases, aligned versions were slightly slower.

This suggests that under the tested conditions and with current compiler optimizations, memory alignment did not provide a measurable performance improvement. The benefits of alignment may be more apparent in SIMD-heavy or low-level hardware-optimized code, or with much larger datasets.

## 5. Inlining

Size: 3

Matrix-Vector-Column | Avg Time: 3.748e-05 ms | Std Dev: 2.3971e-05 ms

Matrix-Vector-Column\_Inlined | Avg Time: 4.332e-05 ms | Std Dev: 1.43853e-05 ms

Size: 20

Matrix-Vector-Column | Avg Time: 0.00101838 ms | Std Dev: 0.000116786 ms

Matrix-Vector-Column\_Inlined | Avg Time: 0.0010508 ms | Std Dev: 0.000131256 ms

Size: 100

Matrix-Vector-Column | Avg Time: 0.0339658 ms | Std Dev: 0.00346175 ms

Matrix-Vector-Column\_Inlined | Avg Time: 0.0333408 ms | Std Dev: 0.00240394 ms

The results show that using the inline keyword has minimal impact on the performance of the matrix-vector column-major multiplication. For all tested sizes (3, 20, and 100), the inlined version performs nearly the same as the non-inlined version, with only very slight differences in average execution time and standard deviation.

Inlining is generally beneficial for small, frequently called functions where the function call overhead is relatively significant compared to the function body. However, in this case, the matrix-vector multiplication is already memory-bound, and the function body is relatively large compared to the cost of a function call. As a result, inlining does not lead to noticeable performance gains.

Modern compilers also automatically inline functions during optimization when beneficial. Therefore, manually adding the inline keyword may have little or no effect, especially when compiler optimizations (e.g., -O3) are enabled.

## with aggressive compiler optimizations - -O3:

Size: 3

Matrix-Vector-Row | Avg Time: 8.328e-05 ms | Std Dev: 2.35856e-05 ms

Matrix-Vector-Column | Avg Time: 8.24e-05 ms | Std Dev: 2.27816e-05 ms

Size: 3x3

Matrix-Matrix-Naive | Avg Time: 0.00019592 ms | Std Dev: 3.14152e-05 ms

Matrix-Matrix-Transposed | Avg Time: 0.00019316 ms | Std Dev: 2.74061e-05 ms

Size: 20

Matrix-Vector-Row | Avg Time: 0.00206078 ms | Std Dev: 7.22887e-05 ms

Matrix-Vector-Column | Avg Time: 0.0020267 ms | Std Dev: 7.97735e-05 ms

Size: 20x20

Matrix-Matrix-Naive | Avg Time: 0.042484 ms | Std Dev: 0.00200875 ms

Matrix-Matrix-Transposed | Avg Time: 0.0427275 ms | Std Dev: 0.00261532 ms

Size: 100

Matrix-Vector-Row | Avg Time: 0.0476383 ms | Std Dev: 0.00578055 ms

Matrix-Vector-Column | Avg Time: 0.04732 ms | Std Dev: 0.00595459 ms

Size: 100x100

Matrix-Matrix-Naive | Avg Time: 3.41364 ms | Std Dev: 0.144528 ms

Matrix-Matrix-Transposed | Avg Time: 3.4067 ms | Std Dev: 0.132095 ms

## without aggressive compiler optimizations - -O0:

Size: 3

Matrix-Vector-Row | Avg Time: 8.01e-05 ms | Std Dev: 2.48397e-05 ms

Matrix-Vector-Column | Avg Time: 9.25e-05 ms | Std Dev: 2.39502e-05 ms

Size: 3x3

Matrix-Matrix-Naive | Avg Time: 0.00018424 ms | Std Dev: 2.76879e-05 ms

Matrix-Matrix-Transposed | Avg Time: 0.00017916 ms | Std Dev: 2.23842e-05 ms

Size: 20

Matrix-Vector-Row | Avg Time: 0.00221586 ms | Std Dev: 0.000907416 ms

Matrix-Vector-Column | Avg Time: 0.00220914 ms | Std Dev: 0.000973124 ms

Size: 20x20

Matrix-Matrix-Naive | Avg Time: 0.0390959 ms | Std Dev: 0.00811997 ms

Matrix-Matrix-Transposed | Avg Time: 0.0382925 ms | Std Dev: 0.00705416 ms

Size: 100

Matrix-Vector-Row | Avg Time: 0.0491375 ms | Std Dev: 0.00783763 ms

Matrix-Vector-Column | Avg Time: 0.0486659 ms | Std Dev: 0.00699852 ms

Size: 100x100

Matrix-Matrix-Naive | Avg Time: 3.4227 ms | Std Dev: 0.15749 ms

Matrix-Matrix-Transposed | Avg Time: 3.41261 ms | Std Dev: 0.124625 ms

Compiler optimizations had a significant effect on performance. Across all tested sizes and operations, the -O3 version consistently outperformed the -O0 version, with lower average execution times and smaller standard deviations.

The performance difference is most apparent for larger matrices (e.g., size 100 and 100x100), where -O3 reduces execution time by a noticeable margin. This is due to the compiler applying optimizations such as loop unrolling, instruction reordering, inlining, and better register usage under -O3. In contrast, -O0 disables most optimizations, resulting in more function call overhead, redundant memory accesses, and less efficient use of CPU resources.

## 6. Profiling

It can be referenced in the analysis.txt and README.md file. Here are some slices of reports.

Flat profile:

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self us/call	total us/call	name
54.62	0.12	0.12	450	267.03	267.03	multiply_mm_transposed_b(double const*, int, int, double const*, int, int, double*)
27.31	0.18	0.06	450	133.51	133.51	multiply_mm_naive(double const*, int, int, double const*, int, int, double*)
11.38	0.21	0.03	3649300	0.01	0.01	int std::uniform_int_distribution<int>::operator()<std::mersenne_twister_engine<unsi
6.83	0.22	0.02	4175900	0.00	0.00	std::mersenne_twister_engine<unsigned long, 32uL, 624uL, 397uL, 31uL, 2567483615uL,
0.00	0.22	0.00	459	0.00	0.00	std::__cxx11::basic_string<char, std::char_traits<char>, std::allocator<char> >::com
0.00	0.22	0.00	168	0.00	0.00	void std::vector<double, std::allocator<double> >::_M_realloc_insert<double const&>(<
0.00	0.22	0.00	150	0.00	0.00	multiply_mv_col_major(double const*, int, int, double const*, double*)
0.00	0.22	0.00	150	0.00	0.00	multiply_mv_row_major(double const*, int, int, double const*, double*)
0.00	0.22	0.00	12	0.00	0.00	void std::__cxx11::basic_string<char, std::char_traits<char>, std::allocator<char> >
0.00	0.22	0.00	1	0.00	0.00	_GLOBAL__sub_I_Z21multiply_mv_row_majorPKdiiS0_Pd
0.00	0.22	0.00	1	0.00	0.00	_GLOBAL__sub_I_main

Call graph (explanation follows)

granularity: each sample hit covers 2 byte(s) for 4.54% of 0.22 seconds

index	% time	self	children	called	name
					<spontaneous>
[1]	67.5	0.00	0.15		benchmark(int, int, std::__cxx11::basic_string<char, std::char_traits<char>, std::allocator<char> >) [1]
		0.08	0.00	300/450	multiply_mm_transposed_b(double const*, int, int, double const*, int, int, double*) [2]
		0.04	0.00	300/450	multiply_mm_naive(double const*, int, int, double const*, int, int, double*) [4]
		0.02	0.01	2608400/3649300	int std::uniform_int_distribution<int>::operator()<std::mersenne_twister_engine<unsigned long, 32uL,
		0.00	0.00	459/459	std::__cxx11::basic_string<char, std::char_traits<char>, std::allocator<char> >::compare(char const*)
		0.00	0.00	150/150	multiply_mv_row_major(double const*, int, int, double const*, double*) [17]
		0.00	0.00	150/150	multiply_mv_col_major(double const*, int, int, double const*, double*) [16]
		0.00	0.00	126/168	void std::vector<double, std::allocator<double> >::_M_realloc_insert<double const&>(__gnu_cxx::__norm
		0.00	0.00	9/12	void std::__cxx11::basic_string<char, std::char_traits<char>, std::allocator<char> >::_M_construct<cf
		0.04	0.00	150/450	benchmark_aligned(int, int) [3]
		0.08	0.00	300/450	benchmark(int, int, std::__cxx11::basic_string<char, std::char_traits<char>, std::allocator<char> >)
[2]	54.5	0.12	0.00	450	multiply_mm_transposed_b(double const*, int, int, double const*, int, int, double*) [2]
					<spontaneous>
[3]	32.5	0.00	0.07		benchmark_aligned(int, int) [3]
		0.04	0.00	150/450	multiply_mm_transposed_b(double const*, int, int, double const*, int, int, double*) [2]
		0.02	0.00	150/450	multiply_mm_naive(double const*, int, int, double const*, int, int, double*) [4]
		0.01	0.00	1040900/3649300	int std::uniform_int_distribution<int>::operator()<std::mersenne_twister_engine<unsigned long, 32uL,
		0.00	0.00	42/168	void std::vector<double, std::allocator<double> >::_M_realloc_insert<double const&>(__gnu_cxx::__norm
		0.00	0.00	3/12	void std::__cxx11::basic_string<char, std::char_traits<char>, std::allocator<char> >::_M_construct<cf

The profiling was conducted on both the naive and transposed matrix multiplication implementations. According to the gprof results, the transposed version `multiply_mm_transposed_b` consumed the most time, accounting for 54.62% of total execution. The naive version `multiply_mm_naive` followed with 27.31%.

This shows that matrix multiplication is the dominant workload in the benchmark. Random number generation and vector resizing functions took minor portions of the runtime, and matrix-vector multiplication functions had negligible impact.

The profiler output confirms that the transposed implementation, despite taking more total time, is being called more heavily and likely handles larger data. The cache-friendly access pattern in the transposed version improves performance per operation. The call graph also indicates that the benchmark function delegates most work to these two implementations, which aligns with our understanding of their roles in the overall computation.

## 7. Optimization Strategies

We applied loop reordering to the baseline matrix-matrix multiplication function `multiply_mm_naive` to improve cache performance. In the original implementation, the innermost loop iterates over `k`, which multiplies `matrixA[i * colsA + k]` and `matrixB[k * colsB + j]`. However, accessing `matrixB[k * colsB + j]` results in non-contiguous memory access due to column-wise traversal in a row-major layout.

We reordered the loops to place the `j` loop inside and the `k` loop in the middle, from

```
for (int i = 0; i < rowsA; ++i)
    for (int j = 0; j < colsB; ++j)
        for (int k = 0; k < colsA; ++k)
            result[i * colsB + j] += matrixA[i * colsA + k] * matrixB[k * colsB + j];
```

to

```
for (int i = 0; i < rowsA; ++i)
    for (int k = 0; k < colsA; ++k) {
        double a_ik = matrixA[i * colsA + k];
        for (int j = 0; j < colsB; ++j)
            result[i * colsB + j] += a_ik * matrixB[k * colsB + j];
    }
```

This change ensures that the inner loop over `j` accesses `matrixB[k * colsB + j]` sequentially in memory (row-wise in a row-major layout), while accessing to `matrixA` is still row-wise. Therefore, It improves spatial locality for both `matrixB` and `result`.